

STAT167 Lab #4 - Spring 2025

Ethan Choi

2025/4/25

Contents

Discussion/Lab #4 instructions	1
Lecture Review - data visualization with ggplot2	2
Load the <code>tidyverse</code> package	2
The <code>mpg</code> data set	2
The complete graphing template in <code>ggplot2</code>	3
Example: <code>geom_boxplot()</code> for one variable	3
Exercise #1	4
Example: <code>geom_boxplot()</code> for two variables	4
Exercise #2	7
Exercise #3	8
Lecture Review - data transformation with dplyr	11
The <code>nycflights13::flights</code> data set	11
<code>filter()</code> - choose rows by their column values	11
Exercise #4	13

Discussion/Lab #4 instructions

This week, we will review some `ggplot2` examples for data visualization and `dplyr` examples for data transformation/manipulation.

- First, download the `rmd` file from Canvas.
- Open this `rmd` file in RStudio and click **Knit -> Knit to PDF** to render it to PDF format. You need to have **LaTeX** installed on the computer to render it to PDF format. If not, you can also render it to HTML format.
- Read this `rmd` file and the rendered `pdf/html` file side-by-side, to see how this document was generated!

- Be sure to play with this document! Change it. Break it. Fix it. The best way to learn R Markdown (or really almost anything) is to try, fail, then find out what you did wrong.
- Read over the `dplyr` example code and the output. If you have any questions about certain functions or parameters, it is the time to ask!
- There are some exercises through out this document. Replace **INSERT_YOUR_ANSWER** with your own answers. Knit the file, and check your results.

Please comment your R code thoroughly, and follow the R coding style guideline (<https://google.github.io/styleguide/Rguide.xml>). Partial credit will be deducted for insufficient commenting or poor coding styles.

Lab submission guideline

- After you completed all exercises, save your file to `FirstnameLastname-SID-lab4.rmd` and save the rendered pdf file to `FirstnameLastname-SID-lab4.pdf`. If you can not knit it to pdf, knit it to html first and then print/save it to pdf format.
- Submit **BOTH** your source `rmd` file and the knitted pdf file to **GradeScope**. Do NOT create a zip file.
- You can submit multiple times, you last submission will be graded.

Lecture Review - data visualization with ggplot2

Load the tidyverse package

```
# install the tidyverse package first if you have not done it yet.
#install.packages("tidyverse") # you can comment out this line after you have installed `tidyverse`

library(tidyverse)
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v readr      2.1.5
## v forcats    1.0.0      v stringr   1.5.1
## v ggplot2    3.5.1      v tibble    3.2.1
## v lubridate  1.9.4      v tidyr     1.3.1
## v purrr      1.0.4
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

The mpg data set

The `mpg` data set contains fuel economy data 1999 - 2008 for 38 popular car models. <https://ggplot2.tidyverse.org/reference/mpg.html>

```
?mpg
## starting httpd help server ... done
glimpse(mpg) # get a glimpse of the mpg data
## Rows: 234
## Columns: 11
## $ manufacturer <chr> "audi", "audi", "audi", "audi", "audi", "audi", "audi", "~
## $ model          <chr> "a4", "a4", "a4", "a4", "a4", "a4", "a4", "a4 quattro", "~
## $ displ          <dbl> 1.8, 1.8, 2.0, 2.0, 2.8, 2.8, 3.1, 1.8, 1.8, 2.0, 2.0, 2.~
## $ year           <int> 1999, 1999, 2008, 2008, 1999, 1999, 2008, 1999, 1999, 200~
## $ cyl            <int> 4, 4, 4, 4, 6, 6, 6, 4, 4, 4, 6, 6, 6, 6, 6, 6, 8, 8, ~
## $ trans          <chr> "auto(l5)", "manual(m5)", "manual(m6)", "auto(av)", "auto~
## $ drv            <chr> "f", "f", "f", "f", "f", "f", "f", "f", "4", "4", "4", "4", "4~
## $ cty            <int> 18, 21, 20, 21, 16, 18, 18, 18, 16, 20, 19, 15, 17, 17, 1~
## $ hwy            <int> 29, 29, 31, 30, 26, 26, 27, 26, 25, 28, 27, 25, 25, 25, 2~
## $ fl             <chr> "p", "p", "p", "p", "p", "p", "p", "p", "p", "p", "p", "p", "p~
## $ class          <chr> "compact", "compact", "compact", "compact", "compact", "c~
```


The complete graphing template in ggplot2

```
ggplot(data = <DATA>) +
  <GEOM_FUNCTION> (
    mapping = aes(<MAPPINGS>),
    stat = <STAT>,          # optional
    position = <POSITION> # optional
  ) +
  <COORDINATE_FUNCTION> + # optional
  <FACET_FUNCTION> +      # optional
  <SCALE_FUNCTION> +      # optional
  <THEME_FUNCTION>        # optional
```

Example: `geom_boxplot()` for one variable

Let's first take a look at the distribution of highway mileage.

```
ggplot(data = mpg) +
  geom_boxplot(mapping = aes(x = "", y = hwy))
```



Exercise #1

If you add `facet_wrap(~class)` to the above code, what will happen?


INSERT_YOUR_ANSWER It will create separate boxplots of highway mileage for each car class into different individual subplots. ***

Example: `geom_boxplot()` for two variables

How highway mileage varies across car classes?


- Plot the distribution of highway mileage by class.

```
ggplot(data = mpg) +  
  geom_boxplot(mapping = aes(x = class, y = hwy))
```




To make the trend easier to see, we can reorder `class` based on the median value of `hwy`.

```
ggplot(data = mpg) +  
  geom_boxplot(mapping = aes(x = reorder(class, hwy, FUN = median),  
                              y = hwy))
```



When you have long variable names, `geom_boxplot()` will work better if you flip it with `coord_flip()`.

```
ggplot(data = mpg) +  
  geom_boxplot(mapping = aes(x = reorder(class, hwy, FUN = median),  
                             y = hwy)) +  
  xlab("class") +  
  coord_flip()
```



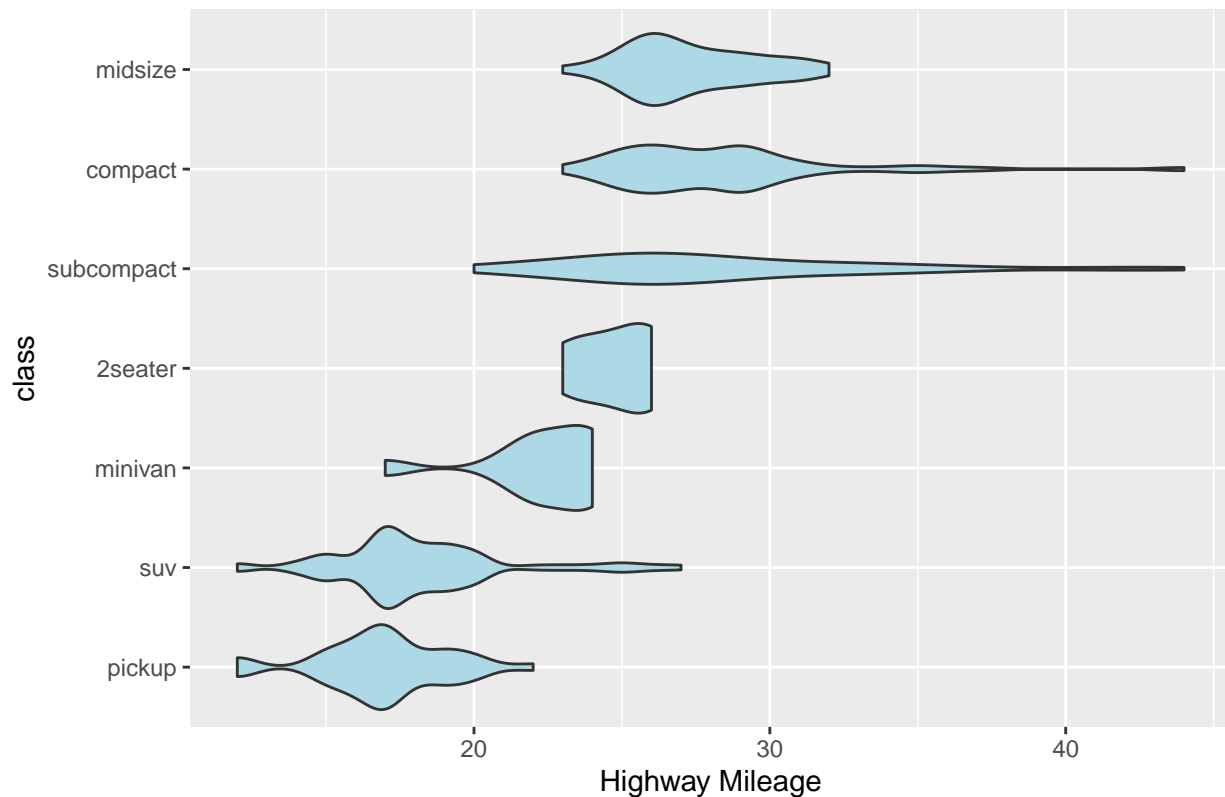
Exercise #2

Can you draw a violin plot to compare the distribution of highway mileage varies across car classes? Order your violins by the median values of highway mileage and flip the plot to horizontal.

INSERT_YOUR_ANSWER

```
ggplot(data = mpg) +
  geom_violin(mapping = aes(x = reorder(class, hwy, FUN = median), y = hwy), fill = "lightblue") +
  coord_flip() +
  xlab("class") +
  ylab("Highway Mileage") +
  ggtitle("Violin Plot of Highway Mileage by Car Class")
```

Violin Plot of Highway Mileage by Car Class




Exercise #3


There are two ways to zoom in a ggplot. Compare the following two plots. What are the differences? Which zooming option do you prefer? Explain your reasons.

INSERT_YOUR_ANSWER The `ylim(c(#, #))` method removes data points outside the range, which can distort/squish the plot and accidentally remove outliers. The `coord_cartesian(ylim = c(#, #))` approach zooms into a specified range without missing any data. I prefer the second method because it keeps the full dataset intact while still focusing the view, which doesn't lead to misinterpretation of data.

```
ggplot(data = mpg) +
  geom_boxplot(mapping = aes(x = class, y = hwy)) +
  ylim(c(20, 30))
## Warning: Removed 100 rows containing non-finite outside the scale range
## (`stat_boxplot()`).
```

```
ggplot(data = mpg) +  
  geom_boxplot(mapping = aes(x = class, y = hwy)) +  
  coord_cartesian(ylim=c(20, 30))
```



Lecture Review - data transformation with dplyr

The nycflights13::flights data set

This data frame contains all 336,776 flights that departed from New York City in 2013. <https://nycflights13.tidyverse.org>

```
# You need to install `nycflights13` first, then you can comment out the following line
#install.packages("nycflights13")
library(nycflights13)
```

```
##?flights # full documentation
glimpse(flights)
## Rows: 336,776
## Columns: 19
## $ year      <int> 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2~
## $ month     <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1~
## $ day       <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1~
## $ dep_time  <int> 517, 533, 542, 544, 554, 554, 555, 557, 557, 558, 558, ~
## $ sched_dep_time <int> 515, 529, 540, 545, 600, 558, 600, 600, 600, 600, 600, ~
## $ dep_delay <dbl> 2, 4, 2, -1, -6, -4, -5, -3, -3, -2, -2, -2, -2, -2, -1~
## $ arr_time  <int> 830, 850, 923, 1004, 812, 740, 913, 709, 838, 753, 849,~
## $ sched_arr_time <int> 819, 830, 850, 1022, 837, 728, 854, 723, 846, 745, 851,~
## $ arr_delay <dbl> 11, 20, 33, -18, -25, 12, 19, -14, -8, 8, -2, -3, 7, -1~
## $ carrier   <chr> "UA", "UA", "AA", "B6", "DL", "UA", "B6", "EV", "B6", "~
## $ flight    <int> 1545, 1714, 1141, 725, 461, 1696, 507, 5708, 79, 301, 4~
## $ tailnum   <chr> "N14228", "N24211", "N619AA", "N804JB", "N668DN", "N394~
## $ origin    <chr> "EWR", "LGA", "JFK", "JFK", "LGA", "EWR", "EWR", "LGA",~
## $ dest      <chr> "IAH", "IAH", "MIA", "BQN", "ATL", "ORD", "FLL", "IAD",~
## $ air_time  <dbl> 227, 227, 160, 183, 116, 150, 158, 53, 140, 138, 149, 1~
## $ distance  <dbl> 1400, 1416, 1089, 1576, 762, 719, 1065, 229, 944, 733, ~
## $ hour      <dbl> 5, 5, 5, 5, 6, 5, 6, 6, 6, 6, 6, 6, 6, 6, 5, 6, 6, 6~
## $ minute    <dbl> 15, 29, 40, 45, 0, 58, 0, 0, 0, 0, 0, 0, 0, 0, 59, 0~
## $ time_hour <dtm> 2013-01-01 05:00:00, 2013-01-01 05:00:00, 2013-01-01 0~
```

filter() - choose rows by their column values

filter() allows you to subset rows based on their column values.

```
# flights that departed on Jan/1/2013
filter(flights, month == 1, day == 1)
## # A tibble: 842 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517             515           2     830             819
## 2  2013     1     1     533             529           4     850             830
## 3  2013     1     1     542             540           2     923             850
## 4  2013     1     1     544             545          -1    1004            1022
## 5  2013     1     1     554             600          -6     812             837
## 6  2013     1     1     554             558          -4     740             728
## 7  2013     1     1     555             600          -5     913             854
## 8  2013     1     1     557             600          -3     709             723
## 9  2013     1     1     557             600          -3     838             846
```

```
## 10 2013      1      1      558      600      -2      753      745
## # i 832 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
# flights that departed in Nov or Dec
filter(flights, month == 11 | month == 12)
## # A tibble: 55,403 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1 2013    11     1       5           2359           6     352           345
## 2 2013    11     1      35           2250          105     123          2356
## 3 2013    11     1     455           500           -5     641           651
## 4 2013    11     1     539           545           -6     856           827
## 5 2013    11     1     542           545           -3     831           855
## 6 2013    11     1     549           600          -11     912           923
## 7 2013    11     1     550           600          -10     705           659
## 8 2013    11     1     554           600           -6     659           701
## 9 2013    11     1     554           600           -6     826           827
## 10 2013    11     1     554           600           -6     749           751
## # i 55,393 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
filter(flights, month %in% c(11, 12))
## # A tibble: 55,403 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1 2013    11     1       5           2359           6     352           345
## 2 2013    11     1      35           2250          105     123          2356
## 3 2013    11     1     455           500           -5     641           651
## 4 2013    11     1     539           545           -6     856           827
## 5 2013    11     1     542           545           -3     831           855
## 6 2013    11     1     549           600          -11     912           923
## 7 2013    11     1     550           600          -10     705           659
## 8 2013    11     1     554           600           -6     659           701
## 9 2013    11     1     554           600           -6     826           827
## 10 2013    11     1     554           600           -6     749           751
## # i 55,393 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
# flights that weren't delayed (on arrival or departure) by more than two hours
filter(flights, !(arr_delay > 120 | dep_delay > 120))
## # A tibble: 316,050 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1 2013     1     1     517           515           2     830           819
## 2 2013     1     1     533           529           4     850           830
## 3 2013     1     1     542           540           2     923           850
```

```
## 4 2013 1 1 544 545 -1 1004 1022
## 5 2013 1 1 554 600 -6 812 837
## 6 2013 1 1 554 558 -4 740 728
## 7 2013 1 1 555 600 -5 913 854
## 8 2013 1 1 557 600 -3 709 723
## 9 2013 1 1 557 600 -3 838 846
## 10 2013 1 1 558 600 -2 753 745
## # i 316,040 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## # tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## # hour <dbl>, minute <dbl>, time_hour <dtm>

filter(flights, arr_delay <= 120, dep_delay <= 120)
## # A tibble: 316,050 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1 2013     1     1     517           515           2     830           819
## 2 2013     1     1     533           529           4     850           830
## 3 2013     1     1     542           540           2     923           850
## 4 2013     1     1     544           545          -1    1004          1022
## 5 2013     1     1     554           600          -6     812           837
## 6 2013     1     1     554           558          -4     740           728
## 7 2013     1     1     555           600          -5     913           854
## 8 2013     1     1     557           600          -3     709           723
## 9 2013     1     1     557           600          -3     838           846
## 10 2013     1     1     558           600          -2     753           745
## # i 316,040 more rows
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## # tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## # hour <dbl>, minute <dbl>, time_hour <dtm>
```

Exercise #4

- (a) Find the flights that departed in the summer (July, August, and September)

INSERT_YOUR_ANSWER

```
summer_flights <- filter(flights, month %in% c(7, 8, 9))
head(summer_flights)
```

```
## # A tibble: 6 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1 2013     7     1     1           2029           212     236           2359
## 2 2013     7     1     2           2359           3     344           344
## 3 2013     7     1    29           2245          104     151           1
## 4 2013     7     1    43           2130          193     322           14
## 5 2013     7     1    44           2150          174     300           100
## 6 2013     7     1    46           2051          235     304           2358
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## # tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## # hour <dbl>, minute <dbl>, time_hour <dtm>
```

(b) Find the flights that flew to Houston (IAH or HOU)

INSERT_YOUR_ANSWER

```
houston_flights <- filter(flights, dest %in% c("IAH", "HOU"))
head(houston_flights)
```

```
## # A tibble: 6 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
## 1  2013     1     1     517           515         2      830           819
## 2  2013     1     1     533           529         4      850           830
## 3  2013     1     1     623           627        -4      933           932
## 4  2013     1     1     728           732        -4     1041          1038
## 5  2013     1     1     739           739         0     1104          1038
## 6  2013     1     1     908           908         0     1228          1219
## # i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```